

# RANDOMIZED BYZANTINE GENERALS

Michael O. Rabin\*

Hebrew University and Harvard University

**Abstract.** We present a randomized solution for the Byzantine Generals Problems. The solution works in the synchronous as well as the asynchronous case and produces Byzantine Agreement within a fixed small expected number of computational rounds, independent of the number  $n$  of processes and the bound  $t$  on the number of faulty processes. The solution uses A. Shamir's method for sharing secrets. It specializes to produce a simple solution for the Distributed Commit problem.

## 0. INTRODUCTION

We present a randomized solution for the well known Byzantine Generals (BG) problem. This problem was introduced by Lamport, Pease and Shostak [LPS80], and was already the subject of numerous interesting studies. In particular, Dolev and Strong [DS81], have shown, extending a result of Lynch and Fischer [LF82], that for  $n$  processes with up to  $t$  faulty ones,  $t+1$  computing phases are necessary for reaching Byzantine Agreement. Another result [FLP82] shows that the BG problem has no solution in the asynchronous case.

Our randomized protocol is certain to achieve Byzantine Agreement, and the expected number of rounds required to do so is four, independent of  $n$  and  $t$ . The total expected number of messages exchanged is  $cnt$ , where  $c$  is a small constant.

Another version of the protocol, employs a fixed number  $R$  of rounds to reach Byzantine agreement, but there is a probability  $2^{-R}$  of error.

The algorithms we use are randomized in the sense of [Ra76], i.e., they employ randomly chosen numbers. Their properties and efficacy do not depend on an assumption of random behavior of the faulty processes.

Like some other solutions, see [DS81], our solution employs authentication of messages by digital signatures. This requires that the processes in the system be supplied in advance by a non-faulty Dealer, with a directory of public-keys.

A salient novel feature is the use of randomly chosen secrets which are shared by the processes in the manner invented by A. Shamir [Sh79]. This also requires that the processes be supplied in advance with certain information by a non-faulty Dealer.

The solution applies, with appropriate modifications, to the synchronized version as well as to the completely asynchronous version of the problem.

A different randomized solution for the BG problem was given, independently, by Ben-Or [BO83]. Ben-Or treats the case of a two-valued initial message, and the processes autonomously toss a coin until a large number of the individual outcomes coincide. This solution requires, in general, a number of rounds and messages exponential in  $n$ .

If we restrict, in our solution, the modes of failure to just breakdowns of processes, we can dispense with authentication and with Shamir's method for sharing a secret. We then get a very simple and robust solution for the Distributed Commit problem.

Thus, in addition to being quite simple, our solution has stronger properties than other existing solutions. It may well lend itself to practical implementation. The author would like to thank D. Dolev and H. Gaifman for many helpful conversations on the topics of this paper.

## 1. BASIC CONCEPTS

Let  $G_i$ ,  $1 \leq i \leq n$ , be processes (the "generals"). Assume that every  $G_i$  can directly exchange messages with every other  $G_j$ .

\* Research supported in part by NSF Grant MCS-8121431.

Every  $G_i$  has a message  $M_i$ , called  $G_i$ 's initial message, and they have to agree on a common value (content) for the message. To this end, each  $G_i$  executes a program  $P_i$  called the Agreement Protocol. The protocols  $P_i$  involves exchanging messages with other processes  $G_j$ , deciding what value to adopt as the common message, and deciding when to stop.

As long as a  $G_i$  computes according to  $P_i$  it is called proper. Once a process  $G_i$  deviates from  $P_i$  it is faulty, and is considered to remain faulty even if, later on, it reverts back to following  $P_i$ .

Each process  $G_i$  has a variable,  $message(i)$ , local to it. This variable holds, at the end of  $G_i$ 's execution of its agreement protocol,  $G_i$ 's version of the message.

**Definition 1.** The processes in a set  $\{G_i : i \in P\}$  are said to reach an agreement about the value of the message if each  $G_i$ ,  $i \in P$ , does stop, and at the end of the execution of their protocols,  $message(i) = message(j)$  for all  $i, j \in P$ .

**Definition 2.** We say that the proper processes have reached Byzantine Agreement if

1. All the proper processes reach agreement.
2. If all proper  $G_i$  have the same initial message  $M_i = M$  then the proper processes agree on  $M$  as the value of the message.

Note that when we talk about Byzantine Agreement, nothing is assumed about the faulty processes.

If all the proper processes have the same initial message then the system will be called proper, otherwise faulty.

**Definition 3.** Let  $n$  and  $t < n$  be integers. A set of protocols (programs)  $P_i$  for the processes  $G_i$ ,  $1 \leq i \leq n$ , is a solution for the Byzantine Generals problem for  $(n, t)$  if the following holds. Let  $S$  be a system of processes  $G_i$  computing according to  $P_i$  and assume that no more than  $t$  processes become faulty. Then whenever the processes have initial messages, the proper  $G_i$  will reach Byzantine Agreement.

A solution for the Byzantine Generals problem will be called a Byzantine Agreement Protocol (BAP).

It should be noted that in a system with up to  $t$  faulty processes, the processes to become faulty are not selected in advance. Rather, we make the worst case assumption that at any time during the

computation, an adversary can select further processes which will become faulty, provided that the total number of faulty processes will not exceed  $t$ .

For randomizing protocols  $P_i$ , which employ randomly chosen numbers in the computation, we have several possible notions of BAP.

We can talk about protocols which achieve Byzantine Agreement with a small probability of error. More precisely, given  $0 < \epsilon$  we say that randomizing protocols  $P_i$ ,  $1 \leq i \leq n$ , are a  $1 - \epsilon$  reliable Byzantine Agreement Protocol for  $(n, t)$ , if for every fixed or randomized behavior of up to  $t$  faulty processes, the proper processes will reach Byzantine Agreement with probability at least  $1 - \epsilon$ .

Another possibility is to demand that for some constant  $C$ , the proper processes achieve Byzantine Agreement within an expected number  $C$  of phases, and this without error. The notion of a computation phase will be explained shortly.

We shall present randomized BAP's of both types.

We shall present the detailed solution for the completely asynchronous case. In this situation we make the assumption that whenever a proper  $G_i$  sends a message to another proper  $G_j$ , the message will eventually reach  $G_j$ . Also, a proper process will execute the instructions of its protocol within a finite time. In fact, these requirements could be incorporated into the definition of a proper process. Throughout the following, we denote by  $t$  the upper bound on the number of processes that may become faulty.

The processes reach agreement on the common message by exchanging information. In such an exchange a process  $G_i$  will request certain information from certain  $k$  other processes. In waiting for the response,  $G_i$  should not wait for more than  $k - t$  replies, because up to  $t$  processes may be faulty and never respond. So that waiting for more than  $k - t$  responses could block a process. These considerations motivate the following

**Definition 4.** A phase in the computation of a proper  $G_i$ , is the time interval between  $G_i$ 's request for information from  $k$  other processes, and receipt of at least  $k - t$  replies. The computation time required by  $G_i$ , once the information is received, is also included in the phase.

The  $k$  in the above definition is determined by the current instruction executed by the protocol  $P_i$ . The actual duration of a phase may be different for

different processes, and may vary from phase to phase for the same  $G_i$ . But our assumptions insure that for a proper process, each phase is finite.

We assume that each process  $G_i$  has a local phase-clock  $p(i)$ , and that  $P_i$  assigns  $p(i) := p(i) + 1$  at the end of each phase.

## 2. AUTHENTICATION

Our solution will use digital signatures for authentication of messages. One implementation, see [DH76], [RSA78], is to have a public directory containing for each participant  $B$  a public key  $K_B$ . When the participant  $B$  needs to authenticate a message  $M$ , he employs a secret key  $D_B$  to compute another message  $\sigma_B(M) = N$ . Every other user can recover  $M$  from  $N$  by use of  $K_B$ , and the fact that  $M$  was so recovered is conclusive proof that it originated with  $B$ . We omit details of the implementation that ensure that nobody other than  $B$  can produce any message of the form  $\sigma_B(M)$ .

For each  $G_i$  we shall denote by  $\sigma_i(M)$  the message  $M$  authenticated by  $G_i$ .

In order to render it impossible for a process  $G_j$  to reuse another processes' old authenticated message  $\sigma_i(M)$ , we employ time-stamping. Each  $G_i$  will incorporate the current reading  $p(i)$  of its phase-clock into any message  $M$  to be authenticated, thus producing  $\sigma_i(M, p(i))$ . By keeping track of  $G_i$ 's most recent time-stamp, a recipient can distinguish between messages newly received from  $G_i$ , and old messages. In the future we shall assume time-stamping, without explicitly mentioning it.

The public-key directory is part of the data in each  $P_i$ , and must be incorporated by a non-faulty "dealer" at the creation of the processes  $G_i$ .

## 3. LOTTERY

Our algorithm will employ a lottery procedure by which the proper  $G_i$ 's can agree on a randomly chosen  $s = 0, 1$ .

In [Sh79] A. Shamir gave a method of sharing a secret  $s$  between  $n$  participants so that every  $k$  or more cooperating participants can reconstruct  $s$ , but no fewer than  $k$  participants can do so. According to Shamir's scheme, a "dealer" who chooses the secret  $s$  to be shared, prepares a sequence  $I_i(s)$ ,  $1 \leq i \leq n$ , and gives  $E_i = I_i(s)$  to  $G_i$ . The  $E_i$ , and Shamir's algorithm, have the property that

from every set  $\{E_{i_1}, \dots, E_{i_k}\}$ , the secret  $s$  can be reconstructed. But no collection of fewer than  $k$  values  $E_i$  determines  $s$ .

For our purposes we assume that the dealer is a non-faulty process  $D$  which prepares in advance the programs for the individual  $G_i$ 's. The dealer  $D$  takes  $k = t + 1$ , where  $t$  is the bound on the number of faulty processes. It randomly and independently chooses a sequence  $s_1, \dots, s_N$  where each  $s_m = 0$  or  $s_m = 1$ . The dealer  $D$  then computes

$$E_i^{(m)} = (I_i(s_m), m) \quad 1 \leq i \leq n, \quad 1 \leq m \leq N.$$

He authenticates all the  $E_i^{(m)}$  with his public-key digital signature  $\sigma_D$  and hands to  $G_i$  the sequence

$$\sigma_D(E_i^{(1)}), \dots, \sigma_D(E_i^{(N)}) .$$

The number  $N$  represents the number of lottery rounds that the processes are expected to play during the duration of the system.

The lottery procedure admits a parameter  $m \leq N$ , so that Lottery( $m$ ) is the  $m$ -th lottery round.

When playing in Lottery( $m$ ),  $G_i$  requests from  $2t$  other processes their  $\sigma_D(E_j^{(m)})$ . Since these messages are signed by  $D$ ,  $G_i$  is certain of their authenticity.

Since there are at most  $t$  faulty processes who may not answer,  $G_i$  will have at least  $k = t + 1$  messages  $(I_{j_1}(s_m), m)$ ,  $\dots, (I_{j_k}(s_m), m)$ . The value  $m$  of the second coordinate will assure him that he is dealing with the  $m$ -th shared secret. Process  $G_i$  will now compute  $s_m$  from the  $I_{j_k}(s_m)$  he has, using Shamir's method.

At the same time  $G_i$  will send his  $\sigma_D(E_i^{(m)})$  to any process  $G$  that requests it.

Thus at the end of the execution of Lottery( $m$ ) by the proper  $G_i$ , at least all the proper processes will share  $s_m$ . Later on we shall explain how the proper  $G_i$  ensure that despite their asynchronous behavior,  $s_m$  is not revealed prematurely.

## 4. THE AGREEMENT PROTOCOL

We shall describe the BAP by a sequence of procedures written informally in Pascal style. Each process has a number of

variables local to it. By the notation  $v(i)$  we mean that  $v(i)$  is local to  $G_i$ , but is global to  $G_i$ 's protocol  $P_i$ .

In particular  $p(i)$  will denote  $G_i$ 's current phase-clock reading. The variable  $message(i)$  will hold at any time  $G_i$ 's current version of the message, which will be one of the initial messages  $M_j$  or else the value "system faulty".

The overall structure of BAP will consist of Polling, where  $G_i$  polls the other processes on their value of the message. Lottery, where the  $G_i$  decide on a common random bit  $s$ . And Decision, where  $G_i$  determines, using  $s$ , whether to adopt the plurality candidate version of the message obtained through Polling, as his current version of the message. This is repeated a fixed number  $R$  of times. The value of  $R$  is determined by the desired reliability, which will turn out to be  $1 - 2^{-R}$ . Thus the BAP will be

Procedure BAP; {for  $G_i$ }

```
begin
  message(i) :=  $M_i$ ; { $G_i$ 's initial message}
  For  $k=1$  to  $k=R$  do { $k$  local to  $P_i$ }
  begin
    Polling; { $k$  is a parameter of Polling}
    Lottery; { $k$  is a parameter of Lottery}
    Decision
  end {For}
end; {BAP}
```

Procedure Polling;

```
begin
  send  $\sigma_i(message(i), k)$  to all; { $\sigma_i$  is
     $G_i$ 's authentication. The  $k$  indicates
    to recipients the version-number.}
  request the  $k$ -th version of  $message(j)$ 
  from all;
  collect incoming  $\sigma_j(message(j), k)$ ,
  one from each responding  $G_j$ , until
   $n - t$  values received; {including own
  value}
  temp(i) := that  $M$  which occurred most
    often among the received
    ( $message(j), k$ ); {the plurality
    candidate}
  count(i) := |{ $j: (message(j), k) =$ 
    ( $temp(i), k$ )}|
    {count of multiplicity of
    most popular  $k$ -th version}
end; {Polling}
```

Remark. Strictly speaking, the collection by  $G_i$  of incoming communications  $\sigma_j(message(j), k)$  is not done just during the  $k$ -th iteration of Polling. Since the system  $S$  is asynchronous, some early  $G_j$  may send its  $k$ -th communication to  $G_i$  while the latter is in an earlier iteration of Polling. Thus  $P_i$  should incorporate a

listening process  $Li$  which runs parallel to all other procedures of  $P_i$ , and stores all incoming communications from the other  $G_j$  as they arrive. The procedures Polling and Lottery collect the appropriate messages from  $Li$ .

We come now to the heart of the randomizing BAP. At the end of Polling every proper  $G_i$  has a value  $temp(i)$  which from his point of view is the plurality candidate for the common message, and a value  $count(i)$  giving the number of times  $G_i$  received the message  $temp(i)$ .

We shall see later that for  $t < n/4$ , if the system is proper, i.e. the initial message of every proper process was  $M$ , then for every proper  $G_i$  we shall have  $temp(i) = M$  and  $n - 2t \leq count(i)$ , after every execution of Polling.

If, however, the system is faulty, then the up to  $t$  faulty  $G_j$ , can cause the different proper  $G_i$  to have different values for  $temp(i)$ , and force  $count(i)$  to almost any value in  $[1, n]$ .

We must provide  $G_i$  with a rule for deciding whether to adopt  $temp(i)$  as the next candidate for value of the common message, or whether the system is faulty. A fixed rule such as  $n/2 \leq count(i)$  is not appropriate, because if the system is faulty then the faulty  $G_j$  can arrange it so that for some proper  $G_i$   $n/2 \leq count(i)$ , and for others  $count(i) < n/2$ . This would foil agreement.

The solution is to base everybody's decision on a comparison  $n/2 \leq count(i)$ , or  $n - 2t \leq count(i)$ , depending on whether a randomly chosen  $s$ , which is not known to the faulty  $G_j$  before the end of Polling by at least one proper process, has value 0 or 1.

Procedures Lottery; {for  $G_i$ }

```
begin
  ask for  $E_j^{(k)}$  from all;
  send  $E_i^{(k)}$  to all;
  wait until  $t$  different values of  $E_j^{(k)}$ 
    have arrived;
  compute  $s_k$  from  $E_i^{(k)}$  and the available
     $E_j^{(k)}$ ;
end; {Lottery}
```

Next we give the procedure for deciding whether to adopt the value  $temp(i)$  resulting from Polling as the new (but not necessarily final) candidate for the value of the message.

Procedure Decision; {for Gi}

```

begin
  s := sk; {sk is the current secret
             for Gi}
  If (s=0 and n/2 ≤ count(i)) or (s=1 and
    n-2t ≤ count(i)) then
    message(i) := temp(i)
  else
    message(i) := "system faulty"
end; {Decision}

```

## 5. CORRECTNESS

**Theorem 1.** Let  $t < n/10$ . All the proper processes will end their execution of BAP. If the system is proper and the initial message of every proper process is  $M$ , then at the end of BAP we shall have  $message(i) = M$  for all proper  $G_i$ .

If the system is faulty then at the end of BAP, with probability at least  $1-2^{-R}$ , all proper processes  $G_i$  will have the same value for  $message(i)$ .

**Proof.** We omit the proof that every proper process will indeed stop. I.e. we prove fully only partial correctness.

Observe that if the proper processes (asynchronously) enter the  $k$ -th iteration of the For statement of BAP in a state where all have the same value  $V$  for their  $k$ -th version of  $message(i)$ , then they will end that iteration with the value  $message(i) = V$ .

This is true because at the end of Polling, every proper  $G_i$  will have  $temp(i) = V$  and  $n-2t \leq count(i)$ . Hence irrespective of whether  $s=0$  or  $s=1$ , Decision will assign  $message(i) := temp(i)$ .

Thus, once the proper  $G_i$  reach the above state, the common value of  $message(i)$  will stay unchanged to the end of the For statement, and the proper  $G_i$ 's will have reached agreement.

In particular, if the system is proper and its initial message is  $M$  for all proper processes, then all proper  $G_i$  will have  $message(i) = M$  at the end of BAP.

Consider the first proper process, say it is  $G_i$ , to have finished its  $k$ -th iteration of Polling. At that time  $G_i$  has accumulated  $n-t$  responses of the form  $\sigma_j(message(j), k)$ . Of these, at least  $n-2t$  are from proper processes. W.l.o.g., let these be the processes  $\{G_1, \dots, G_{n-2t}\} = H$ . Only then does  $G_i$  release its piece  $E_i^{(k)}$  of the  $k$ -th secret  $s_k$ . So the faulty processes can have  $s_k$  only after the  $k$ -th version of  $message(j)$  has been determined by the  $G_j \in H$ .

First assume that there exists a  $V$  so that for at least  $n-4t$  (proper) processes  $G_j \in H$ , the  $k$ -th version is  $message(j) = V$ . Let now  $G_m$  be any proper process. At most  $t$  responses  $\sigma_j(message(j), k)$  from proper processes will not reach  $G_m$  by the time it finished the  $k$ -th iteration of Polling. Therefore, in its  $k$ -th round of Polling it receives responses from at least  $n-5t$  of the above mentioned processes, and all these are of the form  $\sigma_j(V, k)$ . Since  $n/2 < n-5t$ , the value  $V$  is the majority value. Hence  $temp(m) = V$  and  $n/2 < n-5t \leq count(m)$  at the end of the  $k$ -th invocation of Polling, for every proper  $G_m$ .

If now  $s_k = 0$ , which has probability  $1/2$ , independently of the condition on  $H$ , then  $G_m$  will assign  $message(m) := temp(m) = V$  at the end of the  $k$ -th invocation of Decision.

The other possibility is that for every  $V$ ,

$$(1) \quad |\{j : G_j \in H, V = G_j \text{'s } k\text{-th value of } message(j)\}| < n-4t.$$

Let again  $G_m$  be any proper process. Among the  $n-t$  responses  $\sigma_j(message(j), k)$  it collects there are at least  $n-3t$  from  $G_j \in H$ . Thus at most  $2t$  responses from outside of  $H$ . Even if all these responses have the form  $\sigma_\ell(V, k)$  for some common  $V$ , and this  $V$  is the most popular version, (1) will imply that  $count(m) < n-4t+2t = n-2t$  at the end of the  $k$ -th invocation of Polling.

Thus  $count(m) < n-2t$  holds for every proper  $G_m$ . If now  $s_k = 1$  then at the end of the  $k$ -th invocation of Decision all the proper processes will have  $message(m) = \text{"system faulty"}$ . This again will happen with probability  $1/2$ .

Once agreement is reached it persists, so the probability for not reaching agreement in  $R$  rounds is  $2^{-R}$ .

## 6. BOUNDED EXPECTED TIME, ERRORLESS SOLUTION

We turn now to a BAP which ensures a Byzantine agreement without error, and reaches agreement within a small expected number of rounds, independently of  $n$  and  $t$ , as long as  $t < n/10$ .

The key observation is that if a process  $G_i$  finds, during the  $k$ -th iteration of Decision, that  $s_k = 0$ , and  $n-2t \leq count(i)$  and  $s_k = 0$ , then necessarily every other proper  $G_j$  will have  $n/2 < n-4t \leq count(j)$  during its  $k$ -th round of Decision. Hence according to the If statement in Decision, every proper  $G_j$  will assign  $message(j) := temp(j)$  at the end of this invocation of

Decision. By the proof of Theorem 1, all values  $\text{temp}(j)$  for proper  $G_j$  will then be the same. Thus if the above condition holds for some proper  $G_i$ , it constitutes proof for the fact that all proper  $G_j$  will attain the same value for  $\text{message}(j)$  at the end of their  $k$ -th invocation of Decision.

To utilize this fact, we modify Decision. We introduce an integer variable  $k(i)$  which is local to  $G_i$  but global within  $P_i$ .

Procedure Decisionproof; {for  $G_i$ }

```
begin
  s := sk(i); {sk(i) is current shared secret}
  If (s = 0 and n/2 ≤ count(i)) or (s = 1)
    and n - 2t ≤ count(i)) then
    message(i) := temp(i) {= V}
  else
    message(i) := "system faulty";
  If (s = 0 and n - 2t ≤ count(i)) then
    send σi ("agreement reached on V")
    to all Gj;
  k(i) := k(i) + 1 {update the index of the
    secret to be shared}
end; {Decisionproof}
```

Since there will be no fixed number of phases before stopping, we shall need a Boolean-valued variable  $\text{inround}(i)$  to provide  $P_i$  with a halt-signal. The BAP (actually  $P_i$ ) will now read

Procedure ETBAP; {expected time protocol}

```
begin
  inround(i) := true;
  k(i) := 1;
  message(i) := Mi {the initial message
    for Gi}
  repeat
    begin
      Polling;
      Lottery;
      Decisionproof {increases k(i) by 1
        each time}
    end
  end; {ETBAP}
```

**Lemma 2.** Within an expected number of four rounds of iterations of the loop of ETBAP, every proper  $G_j$  will send  $\sigma_j$  ("agreement reached" on  $V$ ) for some  $V$ .

At the time that any proper  $G_j$  sends  $\sigma_j$  ("agreement reached on  $V$ "), it is already determined that all the proper  $G_i$ 's will in fact reach Byzantine Agreement on the value  $V$ .

**Proof.** The procedure Decisionproof is an extension of Decision. Hence it follows,

by the argument in the proof of Theorem 1, that within an expected number of two iterations of the loop, the proper processes will have reached Byzantine Agreement. This means that for some  $V$ , which if the system is proper is its original common message, for all proper  $G_i$  we have  $\text{message}(i) = V$ .

Within an expected number of two additional iterations of the loop, the value  $s_k = 0$  for the shared random secret will appear. At that time (on their asynchronous phase-clocks) all proper  $G_j$ 's will send  $\sigma_j$  ("agreement reached on  $V$ ") during Decisionproof.

As ETBAP stands it does not include a stopping rule. It will not do to stipulate than any  $G_j$  which has sent  $\sigma_j$  ("agreement reached") also assigns  $\text{inround}(j) := \text{false}$  and stops. For the above state may be reached by some proper  $G_j$  before others, and the late proper processes will need cooperation of the early ones in order to reach that state.

We therefore introduce an additional procedure Closefinish which runs simultaneously with ETBAP and provides the stopping rule.

Procedure Closefinish; {for  $G_i$ }

```
begin
  repeat
    If received new σj ("agreement reached
      on V") {Gi's own communication of
      this form is included. V may
      differ for faulty Gi}
    then send σj ("agreement reached V")
      to all
    until t + 1 communications with the same
      V are counted;
    message(i) := V;
    inround(i) := false
  end; {Closefinish}
```

Now, the value  $\text{inround}(i) = \text{false}$  is construed as an interrupt which will cause immediate termination of, and exit from, the repeat statement of ETBAP.

Note that if a proper process has received  $t + 1$  communications  $\sigma_j$  ("agreement reached on  $V$ ") then at least one of these is from a proper  $G_j$ . That proper  $G_j$  actually had a proof that agreement will be reached on the value  $V$ . Thus  $G_i$  correctly stops. The complete proof of correctness is omitted here.

## 7. CONCLUDING REMARKS

We have fully treated in this paper the asynchronous case of the BGP. In the synchronized case it is assumed that the processes  $G_i$  have phase-clocks  $p(i)$ , and that these phases have the same start and end-points for all processes; this is the partially synchronized case. If we further assume that the processes have the same clock-reading at any given time, then we have the fully synchronized case.

The availability of synchronization enables us to simplify our protocol. It now suffices to require  $t < n/4$ . A new issue that arises in the synchronized case is that of coordinated agreement, namely we may want the proper processes to reach final decision about the common message at the same time. The author has an algorithm for enhancing any BA protocol to one that produces coordinated agreement. H. Gaifman produces coordination by use of randomization along the lines of our paper.

Other formulations of the BGP assume a sender  $G$  (who could also be anyone of the  $G_j$ ) who sends each  $G_i$  a message which becomes its initial message  $M_i$ . The case without sender is sometimes called the Byzantine Consensus Problem. The case where there is a sender is readily reduced to the one treated here. All that is additionally required is a Wakeup protocol where any proper process, upon first receiving the sender's message, broadcasts it.

As written here, our solutions require a total of  $O(n^2)$  messages. But this can be readily reduced to  $O(nt)$ .

Finally, the solution extends to cover the situation that the system has to repeatedly reach agreement on new sender's messages. Under appropriate assumptions the solution is resilient. Transient failures of processes in one agreement round, do not affect subsequent rounds.

Details of the above work will be given in the full version of the paper.

## BIBLIOGRAPHY

- [DS81] Dolev, D., Strong, H.R., Polynomial algorithms for multiple processor agreement, IBM Research, San Jose, 1982.
- [D82] Dolev, D., The Byzantine Generals strike again, J. of Algorithms, vol. 3 (1982).

- [LF82] Lynch, N., Fischer, M., A lower bound for the time to assure inter-active consistency, Inf. Proc. Letters, Vol. 4 (1982), pp. 183-186.
- [LSP80] Lamport, L., Shostak, R., and Pease, M., The Byzantine Generals problem, Tech. Report 54, Comp. Sc. Lab., SRI Inter., 1980.
- [BO83] Ben-Or, M., Another advantage of free choice: Completely asynchronous agreement protocols, Abstract.
- [Ra76] Rabin, M.O., Probabilistic algorithms, Algorithms and Complexity, J.F. Traub, Editor, Academic Press (1976), pp. 21-39.
- [FLP82] Fischer, M., Lynch, N., and Paterson, M., Impossibility of distributed consensus with one faulty process, MIT/LCS/TR-282.
- [RSA78] Rivest, R., Shamir, A., and Adleman, L., A method for obtaining digital signatures and public-key cryptosystems, CACM, Vol. 21 (1978), pp. 120-126.
- [Sh79] Shamir, A., How to share a secret, CACM, Vol. 22 (1979), pp. 612-613.
- [DH76] Diffie, W., Hellman, M., New directions in cryptography, IEEE Trans. Inf. Theory, IT-22 (1976), pp. 644-655.